

RAG, Mapped to fetchbot

Christoph Viehoff · OnePointNine LLC · September 2026

fetchbot already ran the retrieval half of a RAG pipeline, with card images instead of text. A maintenance assistant uses the same eight steps with manuals.

#	RAG step	Maintenance assistant	fetchbot
1	Extraction	Pull text from OEM manuals, procedures and work orders	Took in card images from the catalog
2	Chunking	Split documents into paragraph-sized pieces	Sampled regions of each card image; the sampling algorithm decided which parts of the card count confirm
3	Embeddings	Turn each chunk into a vector that captures meaning	Turned each card into an image hash, a compact fingerprint
4	Storage	Vector database: pgvector or Elasticsearch	Elasticsearch held the card hash databases
5	Retrieval	Find the chunks closest to the question	A user's photo was hashed the same way, and Elasticsearch found the nearest matching cards
6	Reranking	Re-score the top hits for true relevance	Chose the best match among close candidates confirm
7	Generation	The LLM writes an answer from the retrieved chunks	Returned the matched card's details from a template: the same answer-from-retrieval idea, without an LLM
8	Guardrails	Flag or refuse answers without good sources; cite the manual	Avoided showing a wrong card when the match was weak confirm

The interview line: "I've built the retrieval half of RAG before. At fetchbot we turned card images into fingerprints, stored them in Elasticsearch and matched a user's photo to the right card in the catalog. A maintenance assistant is the same pipeline: the fingerprints become text embeddings of manuals, and an LLM writes the final answer instead of a template."

Steps marked confirm describe the general approach. Fill in fetchbot's actual method before using them in an interview.